

Social network Graph Link Prediction - Facebook Challenge

In []:

```
#Importing Libraries
# please do go through this python notebook:
import warnings
warnings.filterwarnings("ignore")

import csv
import pandas as pd#pandas to create small dataframes
import datetime #Convert to unix time
import time #Convert to unix time
# if numpy is not installed already : pip3 install numpy
import numpy as np#Do aritmetic operations on arrays
# matplotlib: used to plot graphs
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns#Plots
from matplotlib import rcParams#Size of plots
from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
import math
import pickle
import os
# to install xgboost: pip3 install xgboost
import xgboost as xgb

import warnings
import networkx as nx
import pdb
import pickle
from pandas import HDFStore, DataFrame
from pandas import read_hdf
from scipy.sparse.linalg import svds, eigs
import gc
from tqdm import tqdm
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
```

In []:

```
#reading
from pandas import read_hdf
df_final_train = read_hdf('storage_sample_stage4.h5', 'train_df', mode='r')
df_final_test = read_hdf('storage_sample_stage4.h5', 'test_df', mode='r')
```

In []:

```
df_final_train.describe
```

Out[]:

```
<bound method NDFrame.describe of
jaccard_followers \
0          273084          1505602          1          0
1          832016          1543415          1          0
2         1325247          760242          1          0
3         1368400          1006992          1          0
4          140165          1708748          1          0
...          ...          ...          ...          ...
99997         139353          893843          0          0
99998         910842          704068          0          0
99999         794228          1172755          0          0
100000        949992          1854931          0          0
100001        1642037          1090977          0          0

jaccard_followees  cosine_followers  cosine_followees \
~          ~          ~          ~          ~          ~
```

0	0.000000	0.000000	0.000000
1	0.187135	0.028382	0.343828
2	0.369565	0.156957	0.566038
3	0.000000	0.000000	0.000000
4	0.000000	0.000000	0.000000
...
99997	0.000000	0.000000	0.000000
99998	0.000000	0.000000	0.000000
99999	0.000000	0.000000	0.000000
100000	0.000000	0.000000	0.000000
100001	0.000000	0.000000	0.000000

	num_followers_s	num_followees_s	num_followees_d	...	svd_v_s_3	\
0	6	15	8	...	1.983691e-06	
1	94	61	142	...	-6.236048e-11	
2	28	41	22	...	-2.380564e-19	
3	11	5	7	...	6.058498e-11	
4	1	11	3	...	1.197283e-07	
...	
99997	7	1	10	...	0.000000e+00	
99998	0	4	1	...	1.336987e-12	
99999	0	5	1	...	2.154438e-11	
100000	1	2	0	...	1.620187e-11	
100001	1	14	0	...	7.778089e-09	

	svd_v_s_4	svd_v_s_5	svd_v_s_6	svd_v_d_1	svd_v_d_2	\
0	1.545075e-13	8.108434e-13	1.719702e-14	-1.355368e-12	4.675307e-13	
1	1.345726e-02	3.703479e-12	2.251737e-10	1.245101e-12	-1.636948e-10	
2	-7.021227e-19	1.940403e-19	-3.365389e-19	-1.238370e-18	1.438175e-19	
3	1.514614e-11	1.513483e-12	4.498061e-13	-9.818087e-10	3.454672e-11	
4	1.999809e-14	3.360247e-13	1.407670e-14	0.000000e+00	0.000000e+00	
...	
99997	0.000000e+00	0.000000e+00	0.000000e+00	-3.303718e-12	1.538318e-13	
99998	4.493330e-15	4.528679e-14	5.475207e-18	0.000000e+00	0.000000e+00	
99999	1.566738e-12	2.294564e-13	3.493379e-14	0.000000e+00	0.000000e+00	
100000	1.325874e-15	2.066643e-14	2.662102e-16	-1.142753e-17	5.200344e-17	
100001	8.005625e-13	9.429577e-11	7.386157e-14	-1.657134e-14	2.085059e-14	

	svd_v_d_3	svd_v_d_4	svd_v_d_5	svd_v_d_6
0	1.128591e-06	6.616550e-14	9.771077e-13	4.159752e-14
1	-3.112650e-10	6.738902e-02	2.607801e-11	2.372904e-09
2	-1.852863e-19	-5.901864e-19	1.629341e-19	-2.572452e-19
3	5.213635e-08	9.595823e-13	3.047045e-10	1.246592e-13
4	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
...
99997	1.296745e-06	2.990887e-13	1.589668e-12	7.338551e-14
99998	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
99999	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
100000	3.858875e-15	2.173437e-17	2.241477e-16	3.528355e-20
100001	2.107704e-07	2.652994e-12	2.004727e-14	2.805020e-14

[100002 rows x 54 columns]>

In []:

```
df_final_train.columns
```

Out[]:

```
Index(['source_node', 'destination_node', 'indicator_link',
      'jaccard_followers', 'jaccard_followees', 'cosine_followers',
      'cosine_followees', 'num_followers_s', 'num_followees_s',
      'num_followees_d', 'inter_followers', 'inter_followees', 'adar_index',
      'follows_back', 'same_comp', 'shortest_path', 'weight_in', 'weight_out',
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s',
      'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4',
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1', 'svd_v_s_2',
      'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6', 'svd_v_d_1',
      'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6'],
      dtype='object')
```

In []:

```
df_final_train.head()
```

Out[]:

	source_node	destination_node	indicator_link	jaccard_followers	jaccard_followees	cosine_followers	cosine_followees	nun
0	273084	1505602	1	0	0.000000	0.000000	0.000000	
1	832016	1543415	1	0	0.187135	0.028382	0.343828	
2	1325247	760242	1	0	0.369565	0.156957	0.566038	
3	1368400	1006992	1	0	0.000000	0.000000	0.000000	
4	140165	1708748	1	0	0.000000	0.000000	0.000000	

5 rows x 54 columns

In []:

```
y_train = df_final_train.indicator_link
y_test = df_final_test.indicator_link
```

In []:

```
df_final_train=df_final_train.drop('indicator_link',axis=1)
df_final_test=df_final_test.drop('indicator_link',axis=1)
```

In []:


```
estimators = [10,50,100,250,450]
train_scores = []
test_scores = []
for i in estimators:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=5, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0,
                                min_samples_leaf=52, min_samples_split=120,
                                min_weight_fraction_leaf=0.0, n_estimators=i, n_jobs=-1,random_state=25,verbose=0,warm_start=False)
    clf.fit(df_final_train,y_train)
    train_sc = f1_score(y_train,clf.predict(df_final_train))
    test_sc = f1_score(y_test,clf.predict(df_final_test))
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Estimators = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(estimators,train_scores,label='Train Score')
plt.plot(estimators,test_scores,label='Test Score')
plt.xlabel('Estimators')
plt.ylabel('Score')
plt.title('Estimators vs score at depth of 5')
```

```
Estimators = 10 Train Score 0.9134734626246626 test Score 0.8698418194014671
Estimators = 50 Train Score 0.9202517042475092 test Score 0.9018085650912813
Estimators = 100 Train Score 0.9225489170241707 test Score 0.917456249080863
Estimators = 250 Train Score 0.9229916897506927 test Score 0.9151985855907053
Estimators = 450 Train Score 0.923543372120984 test Score 0.9159283654959068
```

Out[]:

Text(0.5, 1.0, 'Estimators vs score at depth of 5')


Estimators vs score at depth of 5



In []:

```
depths = [3,9,11,15,20,35,50,70,130]
train_scores = []
test_scores = []
for i in depths:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=i, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0,
                                min_samples_leaf=52, min_samples_split=120,
                                min_weight_fraction_leaf=0.0, n_estimators=115, n_jobs=-1, random_state=25,
                                verbose=0, warm_start=False)
    clf.fit(df_final_train, y_train)
    train_sc = f1_score(y_train, clf.predict(df_final_train))
    test_sc = f1_score(y_test, clf.predict(df_final_test))
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('depth = ', i, 'Train Score', train_sc, 'test Score', test_sc)
plt.plot(depths, train_scores, label='Train Score')
plt.plot(depths, test_scores, label='Test Score')
plt.xlabel('Depth')
plt.ylabel('Score')
plt.title('Depth vs score at depth of 5 at estimators = 115')
plt.show()
```

```
depth = 3 Train Score 0.8859874293561507 test Score 0.8611105111979785
depth = 9 Train Score 0.9588823655104346 test Score 0.9238249594813615
depth = 11 Train Score 0.961099653564426 test Score 0.9248900393543363
depth = 15 Train Score 0.963534859215579 test Score 0.9252507411846338
depth = 20 Train Score 0.9635935629198207 test Score 0.9258183960279391
depth = 35 Train Score 0.9637559104653287 test Score 0.9257676846336792
depth = 50 Train Score 0.9637559104653287 test Score 0.9257676846336792
depth = 70 Train Score 0.9637559104653287 test Score 0.9257676846336792
depth = 130 Train Score 0.9637559104653287 test Score 0.9257676846336792
```



In []:

```
from sklearn.metrics import f1_score
```

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

param_dist = {"n_estimators":sp_randint(105,125),
              "max_depth": sp_randint(10,15),
              "min_samples_split": sp_randint(110,190),
              "min_samples_leaf": sp_randint(25,65)}

clf = RandomForestClassifier(random_state=25,n_jobs=-1)

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               n_iter=5,cv=10,scoring='f1',random_state=25,return_train_score=True)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
#print('mean train scores',rf_random.cv_results_['mean_train_score'])

```

mean test scores [0.96213851 0.96183154 0.96037467 0.96207089 0.96285674]

In []:

```
print('mean train scores',rf_random.cv_results_['mean_train_score'])
```

mean train scores [0.96314275 0.96254055 0.96088891 0.9627027 0.96395558]

In []:

```
print(rf_random.best_estimator_)
```

RandomForestClassifier(max_depth=14, min_samples_leaf=28, min_samples_split=111, n_estimators=121, n_jobs=-1, random_state=25)

In []:

```

clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=14, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, #min_impurity_split=None,
                             min_samples_leaf=28, min_samples_split=111,
                             min_weight_fraction_leaf=0.0, n_estimators=121, n_jobs=-1,
                             oob_score=False, random_state=25, verbose=0, warm_start=False)

```

In []:

```

clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)

```

In []:

```

from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))

```

Train f1 score 0.9636780303798496

Test f1 score 0.9257234185733513

In []:

```

%matplotlib inline
from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = ((C.T) / (C.sum(axis=1))).T)

    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

```

```

labels = [0,1]
# representing A in heatmap format
cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=lab
els)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=lab
els)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=lab
els)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()

```


In []:

```


print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)

```

Train confusion_matrix



Test confusion_matrix




In []:

```

from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()

```


```
plt.show()
```



```
In [ ]:
```

```
features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

Feature Importances



Assignments:

1. Add another feature called Preferential Attachment with followers and followees data of vertex. you can check about Preferential Attachment in below link <http://be.amazd.com/link-prediction/>
2. Add feature called svd_dot. you can calculate svd_dot as Dot product between source node svd and destination node svd features. you can read about this in below pdf https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf
3. Tune hyperparameters for XG boost with all these features and check the error metric.

In []:

```
train_graph = nx.read_edgelist('train_pos_after_eda.csv', delimiter=',', create_using=nx.DiGraph(), nodetype=int)
test_graph = nx.read_edgelist('test_pos_after_eda.csv', delimiter=',', create_using=nx.DiGraph(), nodetype=int)
```

In []:

```
from tqdm import tqdm
###https://colab.research.google.com/drive/151M1cMDCjpYl8DKCEICnm4b4lkflLJRJ
def feat(df, viz):
    num_followers_s=[]
    num_followees_s=[]
    num_followers_d=[]
    num_followees_d=[]
    inter_followers=[]
    inter_followees=[]
    for q,ran in tqdm(df.iterrows()):
        try:
            s1=set(viz.predecessors(ran['source_node']))
            s2=set(viz.successors(ran['source_node']))
        except:
            s1=set()
            s2=set()
        try:
            d1=set(viz.predecessors(ran['destination_node']))
            d2=set(viz.successors(ran['destination_node']))
        except:
            d1=set()
            d2=set()
        num_followers_s.append(len(s1))
        num_followees_s.append(len(s2))
        num_followers_d.append(len(d1))
        num_followees_d.append(len(d2))
        inter_followers.append(len(s1.intersection(d1)))
        inter_followees.append(len(s2.intersection(d2)))
    return num_followers_s,num_followees_s,num_followers_d,num_followees_d,inter_followers,inter_followees
```

In []:

```
tr_num_followers_s,tr_num_followees_s,tr_num_followers_d,tr_num_followees_d,tr_inter_followers,tr_inter_followees=feat(df_final_train,train_graph)
test_num_followers_s,test_num_followees_s,test_num_followers_d,test_num_followees_d,test_inter_followers,test_inter_followees=feat(df_final_test,test_graph)
```

```
100002it [00:10, 9809.40it/s]
50002it [00:03, 14064.44it/s]
```

In []:

```
df_final_train['num_followers_d']=tr_num_followers_d
df_final_test['num_followers_d']=test_num_followers_d
```

In []:


```
def preferential_attachment(df_fin):
    ###https://in.coursera.org/lecture/python-social-network-analysis/preferential-attachment-model-abipd
    link_pred_followees=[]
    link_pred_followers=[]
    link_pred_followees=df_fin['num_followees_s']*df_fin['num_followees_d']
    link_pred_followers=df_fin['num_followers_s']*df_fin['num_followers_d']
    return link_pred_followers,link_pred_followees
```

In []:

```
df_final_train[' link_pred_followers'], df_final_train['link_pred_followees']= preferential_attachment(df_final_train)

df_final_test[' link_pred_followers'], df_final_test['link_pred_followees']= preferential_attachment(df_final_test)
```

In []:

```
Adj = nx.adjacency_matrix(train_graph,nodelist=sorted(train_graph.nodes())).asfptype()
```

In []:

```
sadj_col = sorted(train_graph.nodes())
sadj_dict = { val:idx for idx,val in enumerate(sadj_col)}
```

In []:

```
U, s, V = svds(Adj,k=6)
print('U Shape',U.shape)
print('V Shape',V.shape)
print('s Shape',s.shape)
```

```
U Shape (1780722, 6)
V Shape (6, 1780722)
s Shape (6,)
```

In []:

```
def svd(x, S):
    try:
        z = sadj_dict[x]
        return S[z]
    except:
        return [0,0,0,0,0,0]
```

In []:

```
df_final_train.columns
```

Out[]:

```
Index(['source_node', 'destination_node', 'jaccard_followers',
      'jaccard_followees', 'cosine_followers', 'cosine_followees',
      'num_followers_s', 'num_followees_s', 'num_followees_d',
      'inter_followers', 'inter_followees', 'adar_index', 'follows_back',
      'same_comp', 'shortest_path', 'weight_in', 'weight_out', 'weight_f1',
      'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s', 'page_rank_d',
      'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4',
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1', 'svd_v_s_2',
      'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6', 'svd_v_d_1',
      'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6',
      'num_followers_d', ' link_pred_followers', 'link_pred_followees'],
      dtype='object')
```

In []:

```
def dot_prod(df):
```

```
###https://www.geeksforgeeks.org/numpy-dot-python/  
### https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_1  
ink_prediction.pdf
```

```
dot_u=[]  
dot_v=[]  
for index,ran in df.iterrows():  
    s1=svd(ran['source_node'],U)  
    d1=svd(ran['destination_node'],U)  
    s2=svd(ran['source_node'],V.T)  
    d2=svd(ran['destination_node'],V.T)  
    dot_u.append(np.dot(s1,d1))  
    dot_v.append(np.dot(s2,d2))
```

```
return dot_u,dot_v
```

```
(df_final_train['dot_u'],df_final_train['dot_v'])=dot_prod(df_final_train)  
(df_final_test['dot_u'],df_final_test['dot_v'])=dot_prod(df_final_test)
```

```
In [ ]:
```

```
df_final_train.columns
```

```
Out[ ]:
```

```
Index(['source_node', 'destination_node', 'jaccard_followers',  
      'jaccard_followees', 'cosine_followers', 'cosine_followees',  
      'num_followers_s', 'num_followees_s', 'num_followees_d',  
      'inter_followers', 'inter_followees', 'adar_index', 'follows_back',  
      'same_comp', 'shortest_path', 'weight_in', 'weight_out', 'weight_f1',  
      'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s', 'page_rank_d',  
      'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',  
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4',  
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',  
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1', 'svd_v_s_2',  
      'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6', 'svd_v_d_1',  
      'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6',  
      'num_followers_d', 'link_pred_followers', 'link_pred_followees',  
      'dot_u', 'dot_v'],  
      dtype='object')
```

```
In [ ]:
```

```
df_final_train.drop(['source_node', 'destination_node'],axis=1,inplace=True)  
df_final_test.drop(['source_node', 'destination_node'],axis=1,inplace=True)
```

```
In [ ]:
```

```
df_final_train.columns
```

```
Out[ ]:
```

```
Index(['jaccard_followers', 'jaccard_followees', 'cosine_followers',  
      'cosine_followees', 'num_followers_s', 'num_followees_s',  
      'num_followees_d', 'inter_followers', 'inter_followees', 'adar_index',  
      'follows_back', 'same_comp', 'shortest_path', 'weight_in', 'weight_out',  
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s',  
      'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',  
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3', 'svd_u_s_4',  
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',  
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1', 'svd_v_s_2',  
      'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6', 'svd_v_d_1',  
      'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6',  
      'num_followers_d', 'link_pred_followers', 'link_pred_followees',  
      'dot_u', 'dot_v'],  
      dtype='object')
```

```
In [ ]:
```

```
from sklearn.metrics import f1_score  
from sklearn.model_selection import RandomizedSearchCV  
from scipy.stats import randint as sp_randint
```

```
from scipy.stats import uniform
from xgboost import XGBClassifier
```

###<https://www.projectpro.io/recipes/use-xgboost-classifier-and-regressor-in-python>

```
param = {'min_child_weight': [1, 5, 10],
         'max_depth': [15,20,30,40],
         'learning_rate': [0.03,0.05,0.1,0.15,0.2],
         'n_estimators' : [20,50,100,150],
         'subsample': [0.6, 0.8, 1.0],
         'colsample_bytree': [0.6, 0.8, 1.0],
         'gamma': [1, 1.5, 2, 5]}

clf = XGBClassifier(random_state=25,n_jobs=-1)

rnd_src = RandomizedSearchCV(clf, param_distributions=param,n_iter=3,
                             cv=3,scoring='f1',random_state=25,return_train_score=True)

rnd_src.fit(df_final_train,y_train)
print('mean test scores',rnd_src.cv_results_['mean_test_score'])
print('mean train scores',rnd_src.cv_results_['mean_train_score'])
```

```
mean test scores [0.97800923 0.97994223 0.97307308]
mean train scores [0.98788082 0.99413852 0.97810217]
```

In []:

```
clf1=rnd_src.best_estimator_
clf1.fit(df_final_train,y_train)

y_train_pred = clf1.predict(df_final_train)
y_test_pred = clf1.predict(df_final_test)

print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))
```

```
Train f1 score 0.9950570342205324
Test f1 score 0.9311120543293718
```

In []:

```
from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = ((C.T) / (C.sum(axis=1))).T)

    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)

    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)

    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)

    plt.xlabel('Predicted Class')
```


```
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()
```


In []:

```
print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```

Train confusion_matrix




Test confusion_matrix



In []:

```
from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```




In []:

```
features = df_final_train.columns
```

```

features = df_train.columns
importances = clf1.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(20,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()

```



In []:

```

from prettytable import PrettyTable
t = PrettyTable()
t.field_names=["model","train_f1_Score","test_f1_score"]
t.add_row(["XGBoost","0.995","0.931"])
print(t)

```

```

+-----+-----+-----+
| model | train_f1_Score | test_f1_score |
+-----+-----+-----+
| XGBoost | 0.995 | 0.931 |
+-----+-----+-----+

```

Observations:

1. When actual prediction takes place in the xgboost classifier, the f1 score for test is 0.931.
2. The f1 score for train is 0.995.
3. Some of the Best features that were extracted are:

- inter_followers
- follows_back
- cosine_followers
- jaccard_followees

1. Area under ROC curve is coming out to be 0.93